

Exploratory Data Analysis (EDA) Report: Retail Sales Analysis

Exploratory Data Analysis (EDA) is an essential step in data science to understand datasets, discover patterns, identify anomalies, and extract meaningful insights. In this project, we perform an end-to-end Exploratory Data Analysis on a Retail Sales Dataset sourced from Kaggle. The analysis involves data loading, cleaning, transformation, visualization, and interpretation to support business decision-making.

Procedure: Exploratory Data Analysis

1. Understand the Problem and the Data

2. Import and Inspect the Data

3. Handling Missing Values and Duplicates

4. Explore Data Characteristics

5. Perform Data Transformation

6. Visualize Data Relationships

7. Handling Outliers

8. Communicate Findings and Insights

Dataset Name: retail_sales.csv

Source: Kaggle (Public Dataset)

Project Structure

```
EDA_PROJECT/
|—— data/
|   |—— retail_sales5.csv
|—— modules/
|   |—— data_import.py
|   |—— data_cleaning.py
|   |—— transformation.py
|   |—— stats_analysis.py
|   |—— visualization.py
|   |—— modeling.py
|—— outputs/
|   |—— retail_sales_backup.csv
|—— main.py
```

1. Understand the Problem and the Data

The objective of this project is to analyze retail sales data to identify relationships between Quantity, Sales, Profit, and Category. The analysis helps in understanding sales distribution, profitability trends, and category-wise performance.

Variables in the Dataset:

Numerical Variables: Quantity, Sales, Profit

Categorical Variables: Category

2. Import and Inspect the Data

File Name: data_import.py

The dataset is loaded into a Pandas DataFrame using the `read_csv()` function. This step allows efficient access and manipulation of the retail data.

The `load_data()` function performs the following tasks:

Displays the first five rows using `head()`

Displays dataset structure using `info()`

Helps identify data types and missing values

This step ensures familiarity with the dataset before proceeding further.

First Five Rows of Dataset									
	Row ID	Order ID	Order Date	Ship Date	...	Sales	Quantity	Discount	Profit
0	1	CA-2016-152156	11/8/2016	11/11/2016	...	261.9600	2	0.00	41.9136
1	2	CA-2016-152156	11/8/2016	11/11/2016	...	731.9400	3	0.00	219.5820
2	3	CA-2016-138688	6/12/2016	6/16/2016	...	14.6200	2	0.00	6.8714
3	4	US-2015-108966	10/11/2015	10/18/2015	...	957.5775	5	0.45	-383.0310
4	5	US-2015-108966	10/11/2015	10/18/2015	...	22.3680	2	0.20	2.5164

```
[5 rows x 21 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9994 entries, 0 to 9993
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Row ID          9994 non-null   int64  
 1   Order ID        9994 non-null   object  
 2   Order Date      9994 non-null   object  
 3   Ship Date       9994 non-null   object  
 4   Ship Mode       9994 non-null   object  
 5   Customer ID    9994 non-null   object  
 6   Customer Name   9994 non-null   object  
 7   Segment         9994 non-null   object  
 8   Country         9994 non-null   object  
 9   City             9994 non-null   object  
 10  State            9994 non-null   object  
 11  Postal Code    9994 non-null   int64  
 12  Region          9994 non-null   object  
 13  Product ID     9994 non-null   object  
 14  Category        9994 non-null   object  
 15  Sub-Category   9994 non-null   object  
 16  Product Name   9994 non-null   object  
 17  Sales            9994 non-null   float64 
 18  Quantity        9994 non-null   int64  
 19  Discount        9994 non-null   float64 
 20  Profit           9994 non-null   float64 
dtypes: float64(3), int64(3), object(15)
memory usage: 1.6+ MB
None
```

Ln 11, Col 1 Spaces: 4

3. Handling Missing Values and Duplicates

File Name: `data_cleaning.py`

3.1 Detection of Missing Values

Missing values were identified using `isnull().sum()`. Any missing numerical values in Quantity, Sales, or Profit were handled appropriately.

3.2 Mean Imputation

Missing numerical values were replaced with the mean of their respective columns to preserve data distribution.

3.3 Duplicate Removal

Duplicate records were checked and removed using `drop_duplicates()` to maintain data integrity.

3.4 Final Dataset

After cleaning, the dataset became consistent and free from missing values and duplicates. A cleaned backup file was saved as `retail_sales_backup.csv`.

Cleaned Dataset Saved as `retail_sales_backup.csv`

4. Explore Data Characteristics

File Name: stats_analysis.py

Descriptive statistics were calculated to understand the distribution and variability of retail data.

The following statistical measures were computed:

- **Mean** – Average value of the dataset
- **Median** – Middle value when data is arranged in order
- **Mode** – Most frequently occurring value
- **Standard Deviation** – Measure of how spread out the values are
- **Minimum** – Smallest value in the dataset
- **Maximum** – Largest value in the dataset

Sales and Profit showed moderate variability, indicating different performance levels across transactions and categories.

Descriptive Statistics			
	Quantity	Sales	Profit
count	9994.000000	9994.000000	9994.000000
mean	3.789574	229.858001	28.656896
std	2.225110	623.245101	234.260108
min	1.000000	0.444000	-6599.978000
25%	2.000000	17.280000	1.728750
50%	3.000000	54.490000	8.666500
75%	5.000000	209.940000	29.364000
max	14.000000	22638.480000	8399.976000

5. Perform Data Transformation

File Name: transformation.py

To prepare the data for further analysis, numerical features were standardized.

Standardization

Quantity, Sales, and Profit were scaled using StandardScaler

This ensures fair comparison between features with different units

```
Standardized Data
[[ -0.8043034  0.0515104  0.05659251]
 [ -0.35486486  0.80563348  0.81505408]
 [ -0.8043034 -0.34536777 -0.09300169]
 ...
 [-0.8043034  0.04608048 -0.03954647]
 [ 0.09457367 -0.32133108 -0.06547279]
 [ -0.8043034  0.02134419  0.18907752]]
```

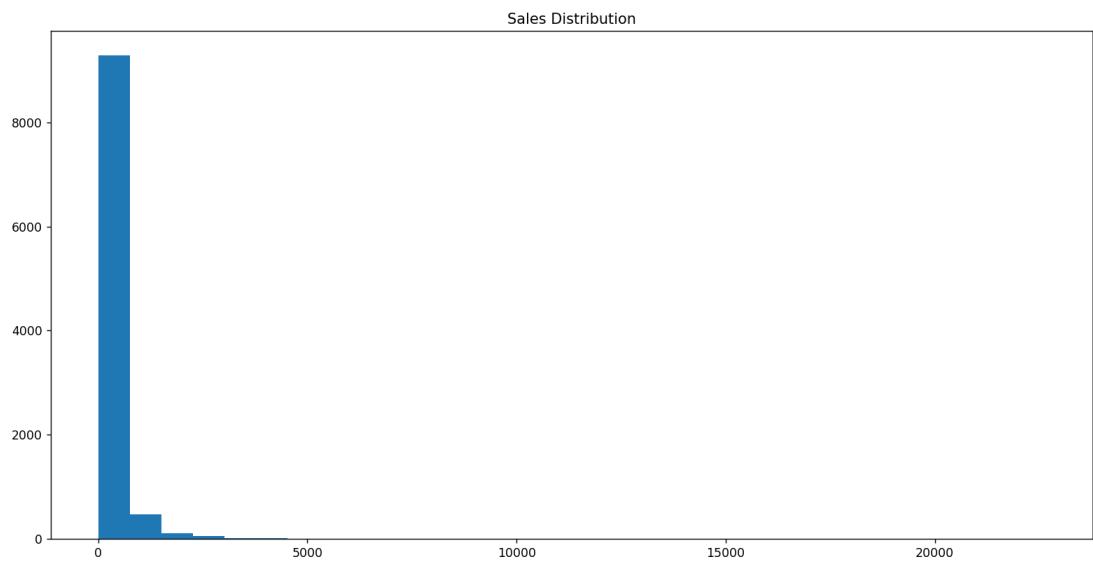
6. Visualize Data Relationships

File Name: visualization.py

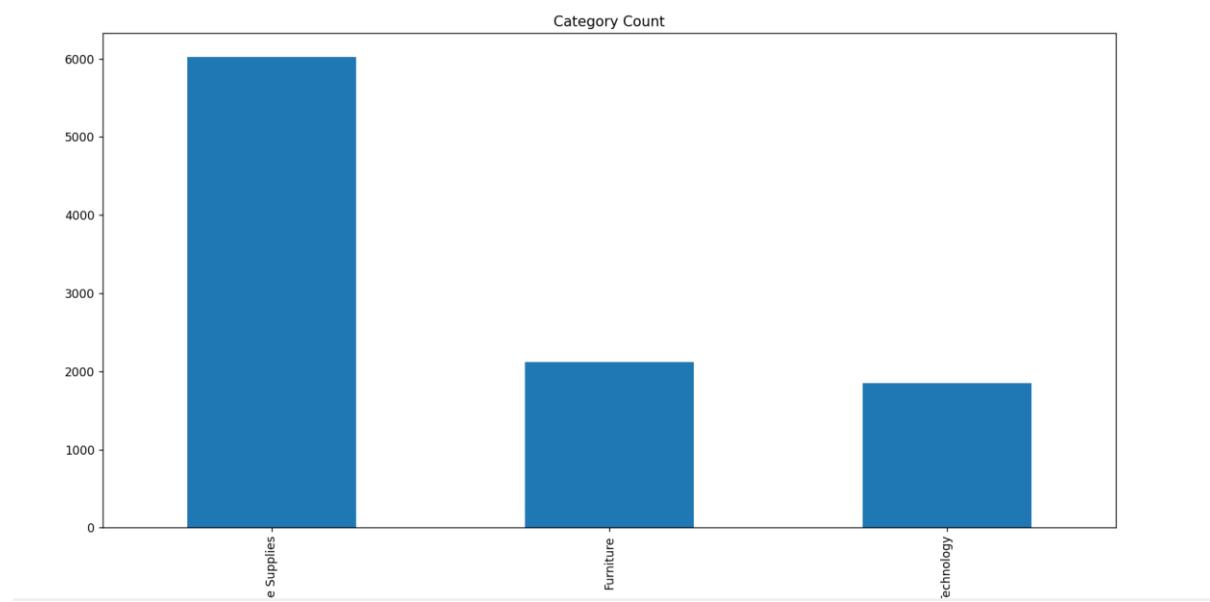
Visualization techniques were used to better understand sales trends and relationships.

Univariate Analysis

Histogram: Distribution of Sales and Profit

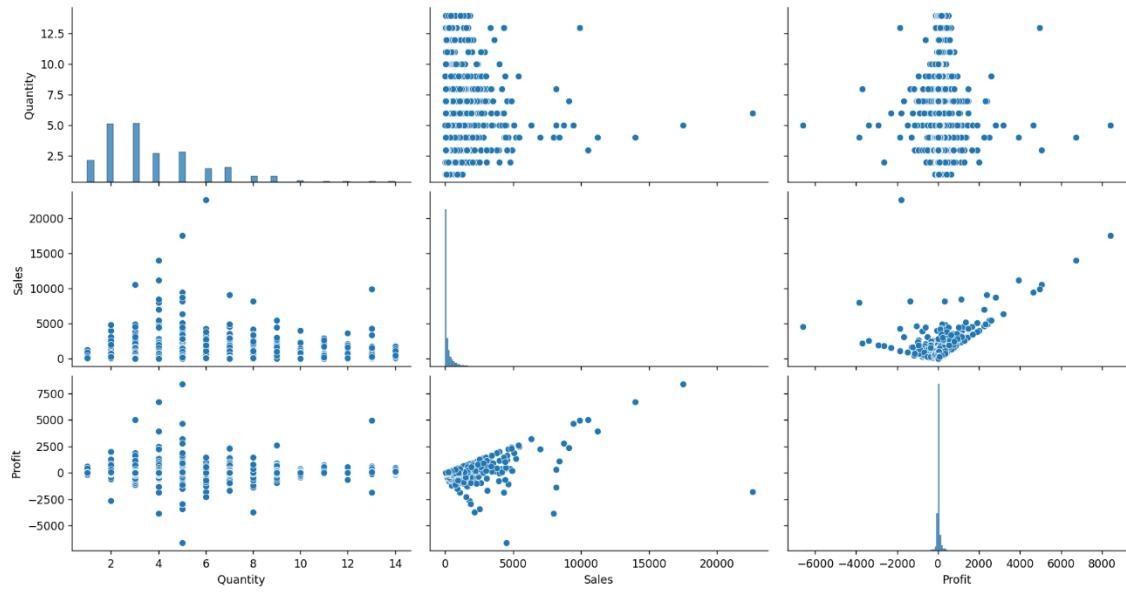


Bar Chart: Category-wise sales count

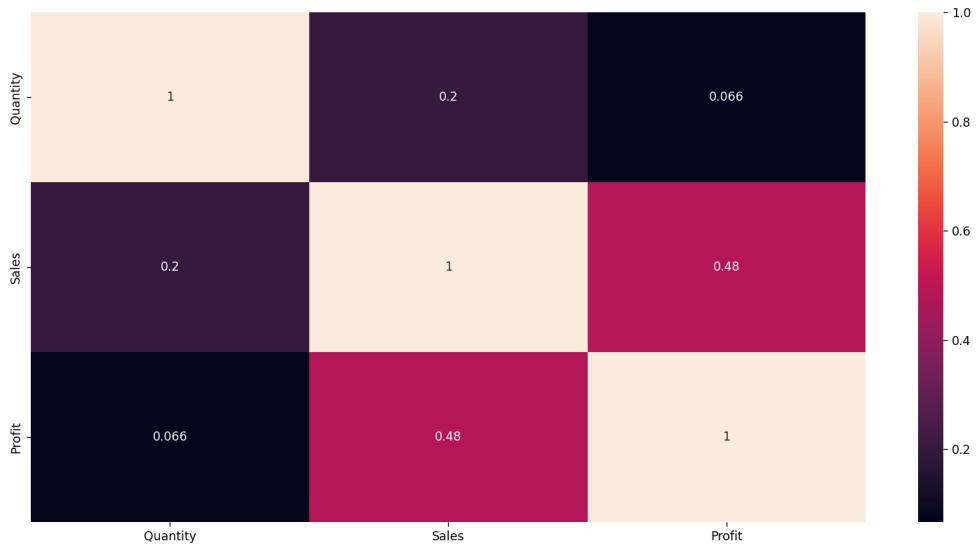


Multivariate Analysis

Pair Plot: Relationship among Quantity, Sales, and Profit



Correlation Heatmap: Strength of relationships between numerical variables



7. Handling Outliers

Outliers were identified using box plots and statistical thresholds. Extreme values in Sales and Profit were reviewed to avoid skewed analysis.

Logic:

Calculate Q1(25th percentile) and Q3(75th percentile).

Define bounds: Lower = Q1 - 1.5 * IQR; Upper = Q3 + 1.5 * IQR

	Quantity	Sales	Profit
count	9994.000000	9994.000000	9994.000000
mean	3.789574	229.858001	28.656896
std	2.225110	623.245101	234.260108
min	1.000000	0.444000	-6599.978000
25%	2.000000	17.280000	1.728750
50%	3.000000	54.490000	8.666500
75%	5.000000	209.940000	29.364000
max	14.000000	22638.480000	8399.976000

8. Communicate Findings and Insights

The final step involved summarizing insights derived from the analysis.

Key Insights

Sales and Profit show a strong positive correlation

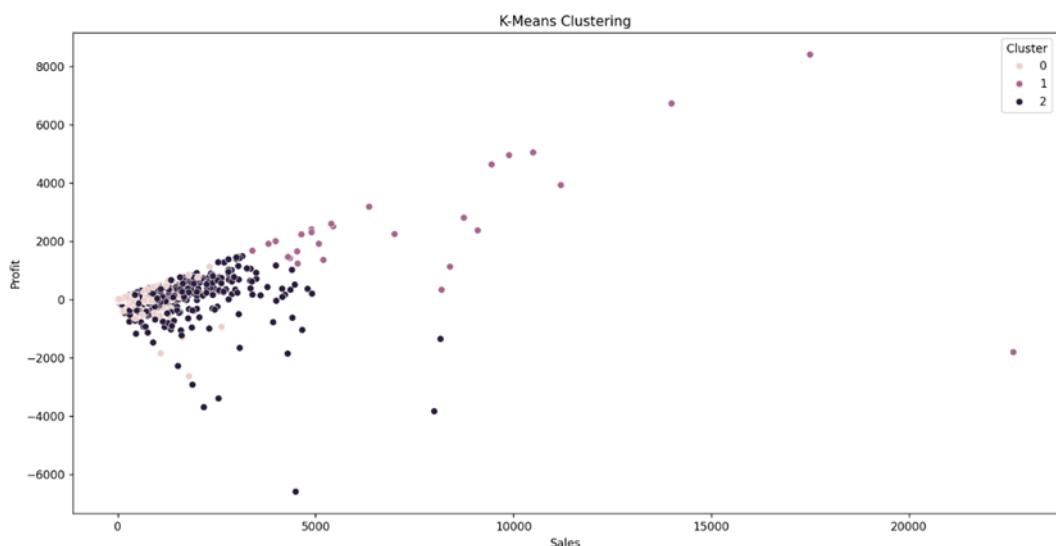
Higher quantity generally leads to increased sales

Category-wise analysis reveals uneven sales distribution

These insights can help businesses optimize pricing strategies, inventory planning, and category management.

8.1 k-means clustering

K-Means clustering is an unsupervised machine learning algorithm used to group data into clusters based on similarity. In this project, K-Means is applied to retail sales data to identify meaningful customer or product groups using numerical features.



CONCLUSION:

This project successfully demonstrates a complete Exploratory Data Analysis workflow on a Retail Sales dataset. By applying data cleaning, statistical analysis, and visualization techniques, meaningful insights were extracted to support data-driven decision-making. The project highlights the importance of EDA in understanding business data and improving retail performance.